

Chapter 1

Introduction

By Prof. Dr.
Amany Sarhan
(2021)

What is distributed system?

- Everywhere in nature we find examples of distribution of activities ???? Give examples!!!!
- Thus, centralizing of activities is not the natural way of handling things.
- Centralization is the opposite of distribution.
- Make definition of centralization !!!!

Motivation towards Distributed systems

- Many applications are inherently distributed like banks, universities,
- Sharing of resources (can be HW like CPU, printers, memory or can be SW like data, database, programs)
- **Parallel** versus **sequential**.
- Time is the key issue in parallel versus sequential comparison.
- If we can process certain tasks at the same time, we are processing in parallel.

Distributed System Definition:

- A distributed system is a collection of **autonomous** hosts that are connected through a computer network.
- Each host executes components and operates a distribution **middleware**.
- Middleware **enables** the components to **coordinate** their activities.
- Users perceive the system as a **single, integrated** computing facility.

Definition of a Distributed system

- A *distributed system* is one in which **components** located at networked computers communicate and **coordinate** their actions only by passing **messages**. This definition leads to the following characteristics of distributed systems:
- Concurrency of components
- Lack of a global 'clock'
- Independent failures of components

Definition of a Distributed system

- There are many other definitions of distributed system, we will concentrate on some of them.

- 1- A distributed system is a collection of independent computers that appear to the users of the system as a single computer***
- 2- A group of connected components that are cooperating to perform a single task (in parallel).
- 3- Interconnected devices that are sharing data and resources.

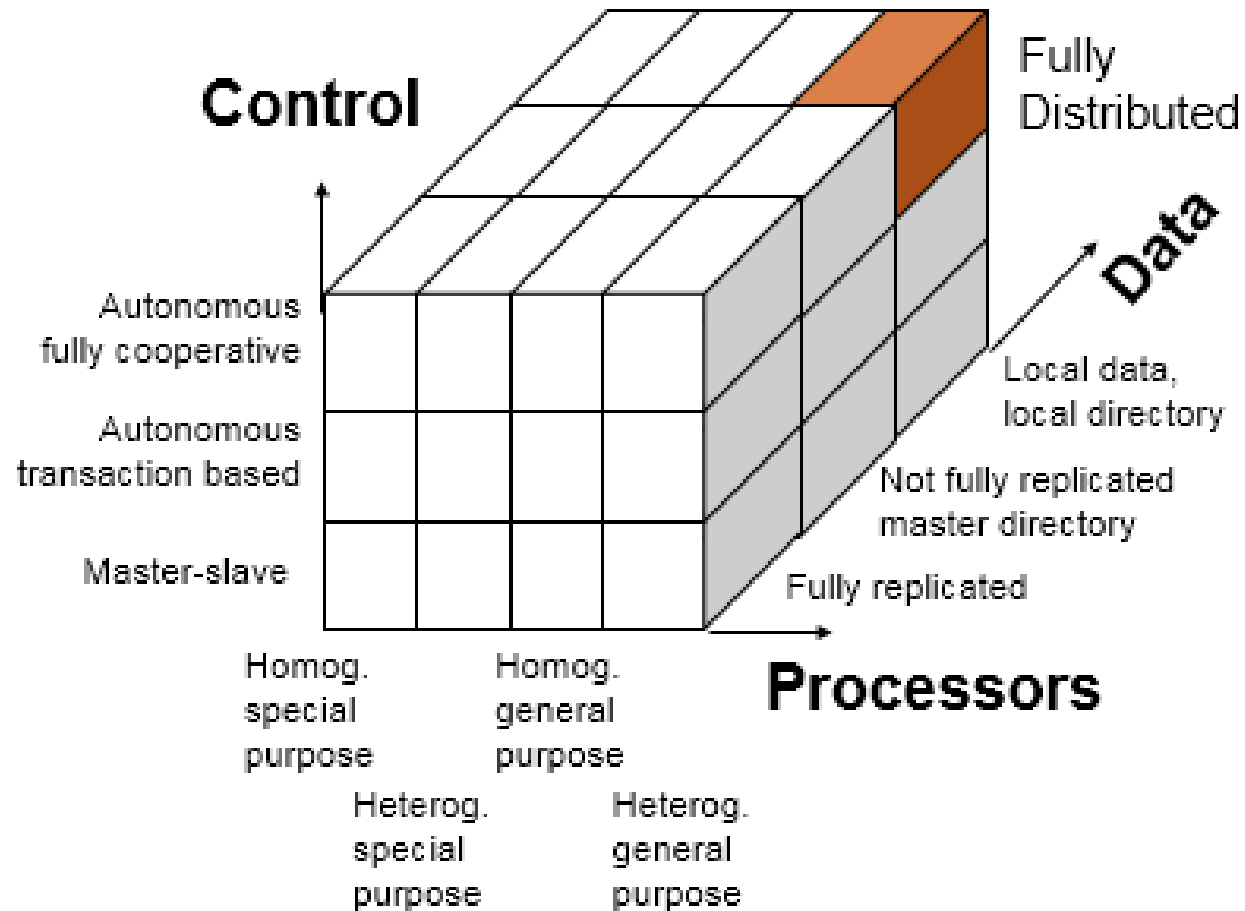
Definition of a Distributed system

4- A number of interconnected autonomous computers that provide services to meet the information processing needs of modern enterprises.

5- A distributed system consists of a collection of autonomous computers linked by a computer network and equipped with distributed system software. This software enables computers to coordinate their activities and to share the resources of the system hardware, software, and data.

➤ Enslow's Definition

➤ Distributed System = Distributed hardware + Distributed control + Distributed data



Definition of a Distributed system

A collection of (probably heterogeneous) automata whose distribution is transparent to the user so that the system appears as one local machine.

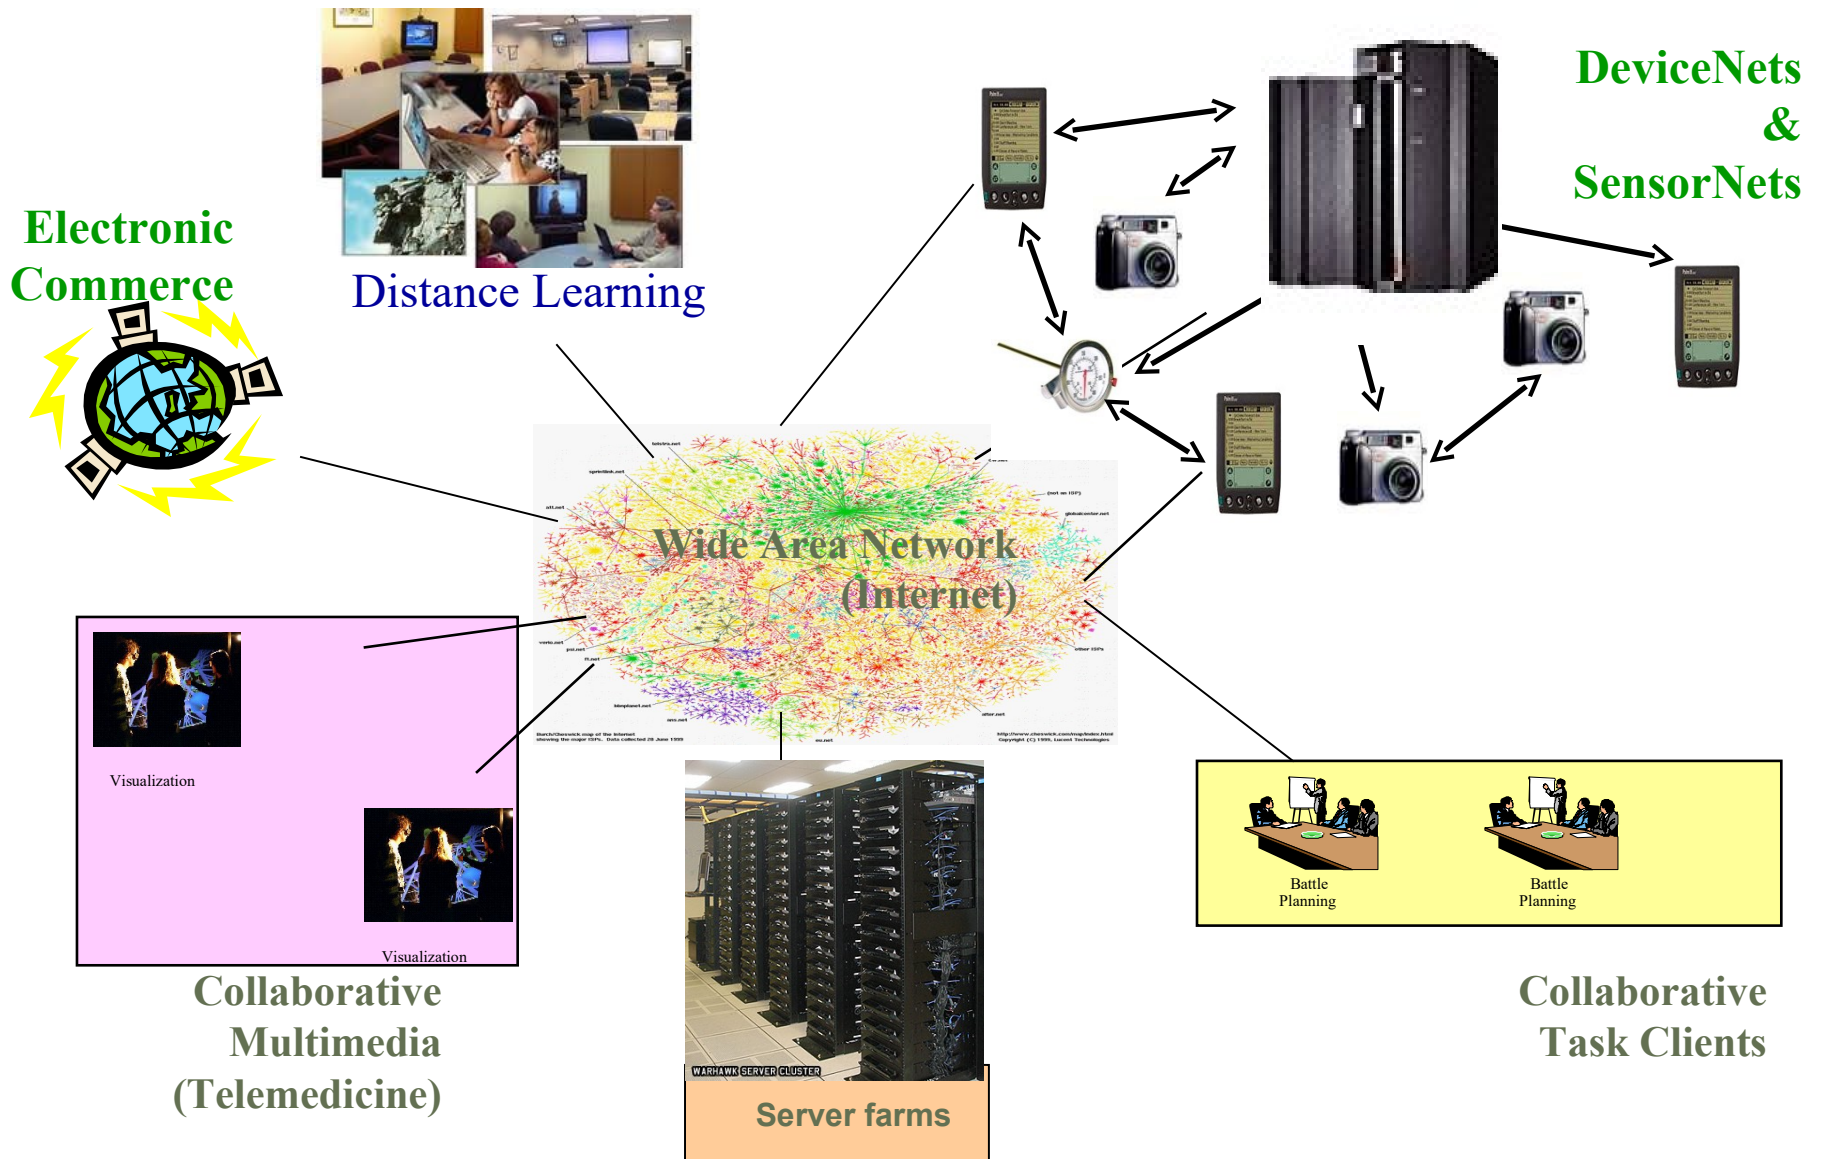
This is in contrast to a network, where the user is aware that there are several machines, and their location, storage replication, load balancing and functionality is not transparent.

Distributed systems usually use some kind of client-server organization.

Distributed Systems application domains connected with networking:

| | |
|---------------------------------------|--|
| Finance and commerce | eCommerce e.g. Amazon and eBay, PayPal, online banking and trading |
| The information society | Web information and search engines, ebooks, Wikipedia; social networking: Facebook and MySpace |
| Creative industries and entertainment | online gaming, music and film in the home, user-generated content, e.g. YouTube, Flickr |
| Healthcare | health informatics, on online patient records, monitoring patients |
| Education | e-learning, virtual learning environments; distance learning |
| Transport and logistics | GPS in route finding systems, map services: Google Maps, Google Earth |
| Science | The Grid as an enabling technology for collaboration between scientists |
| Environmental management | sensor technology to monitor earthquakes, floods or tsunamis |

Next Generation Information Infrastructure



Requirements - Availability, Reliability, Quality-of-Service, Cost-effectiveness, Security

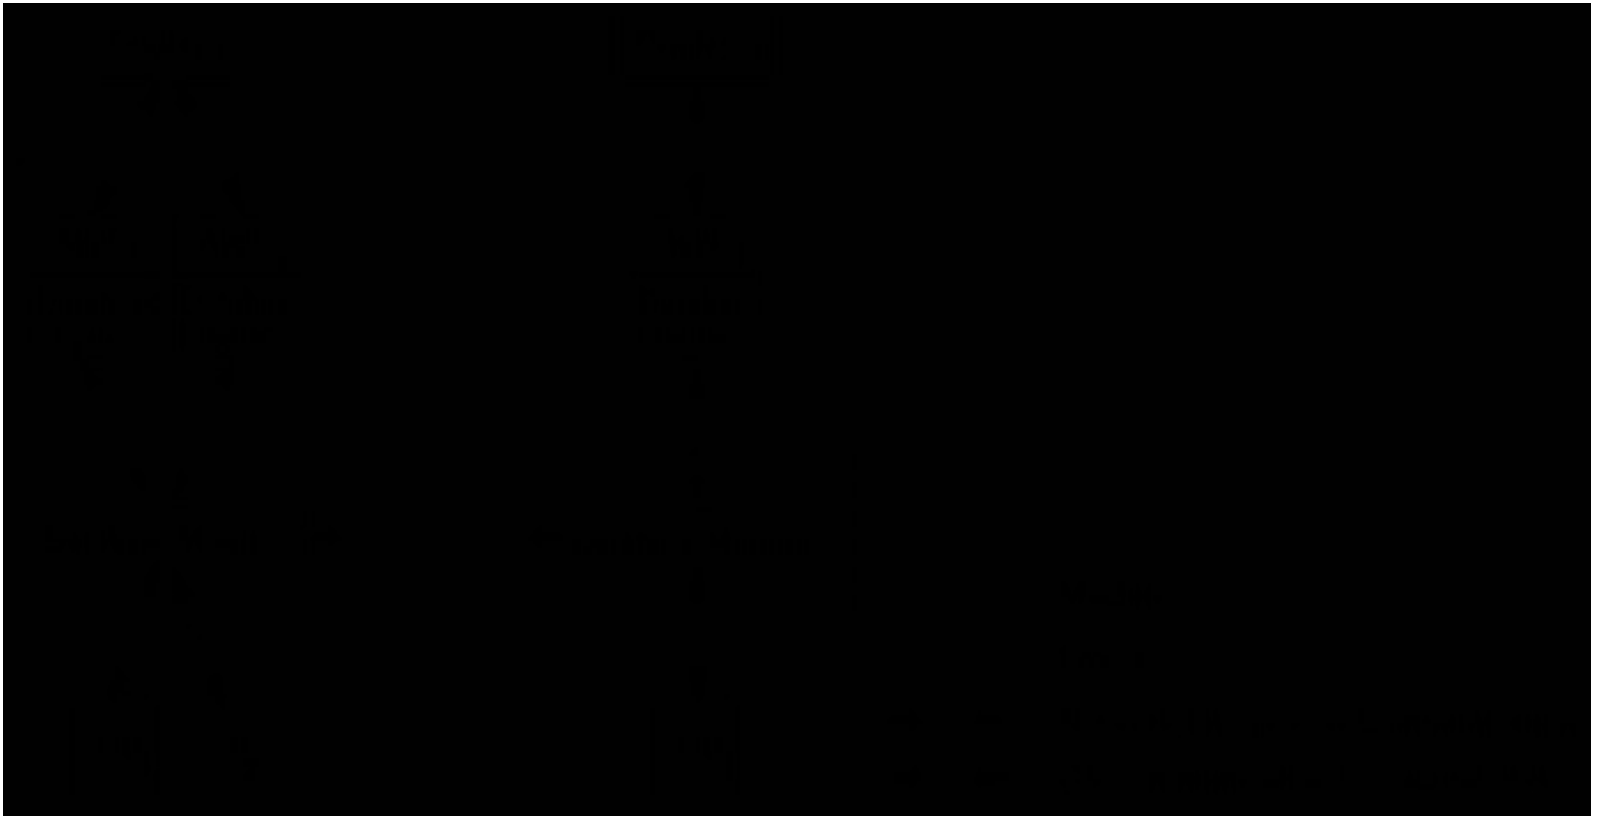
Distributed System Characteristics

- Multiple autonomous components (They possess full control over their parts at all times. The components, however, have to provide interfaces to be able to use each other)
- Components are not shared by all users (components that are used by some only users but are not used by others.)
- Software runs in concurrent processes on different processors
- Multiple points of control (but these are not totally independent. Components have to take into account that they are being used by other components and have to react properly to requests)
- Multiple points of failure (The system may fail because a component of the system has failed or if the network has broken down)

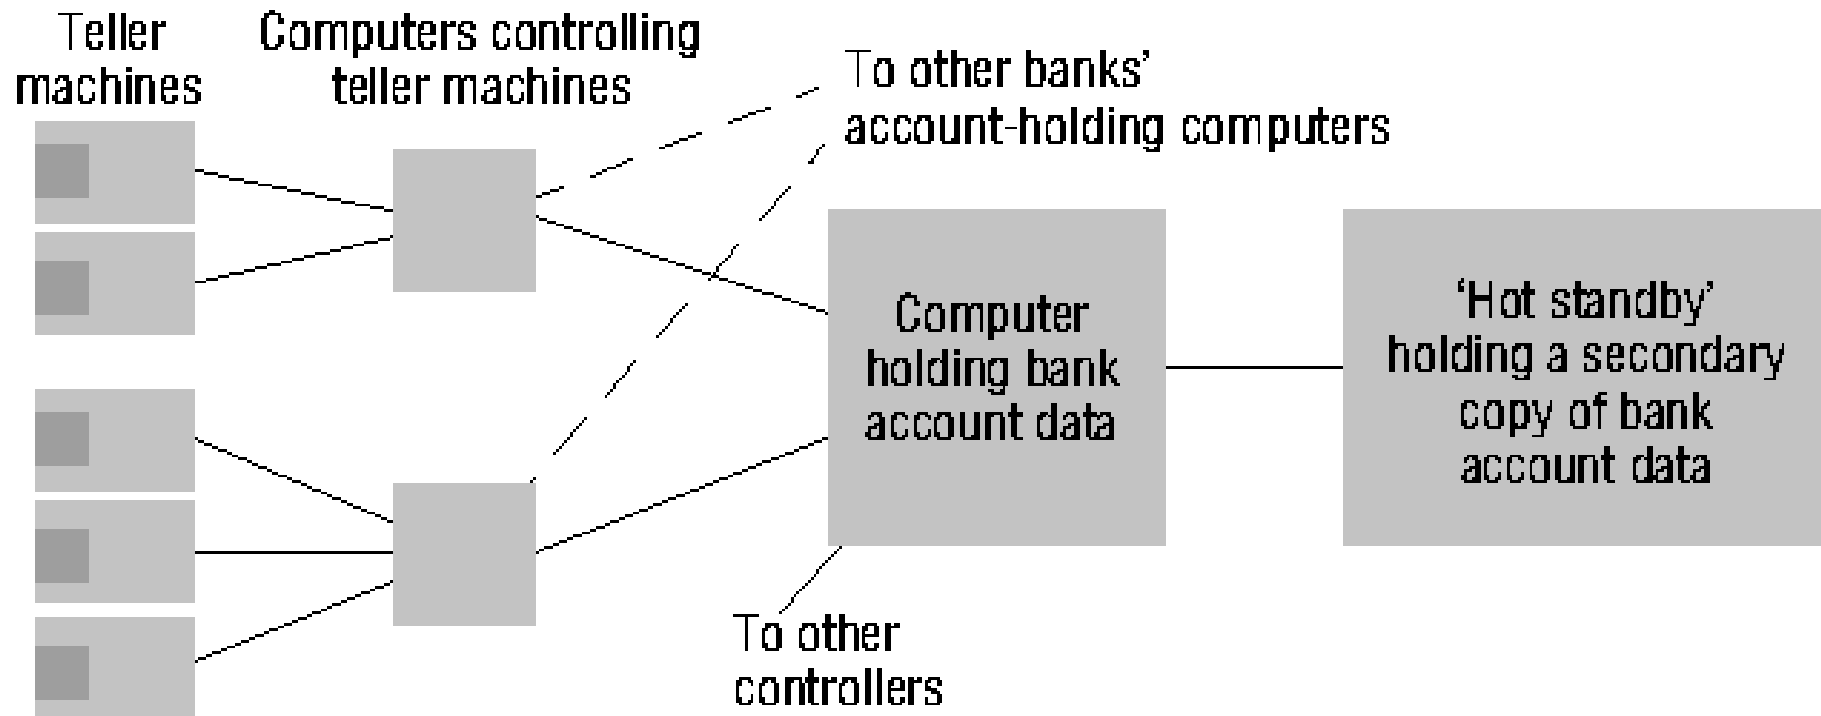
Examples of Distributed Systems

- Database Management System
- Automatic Teller Machine Network
- Internet
- Mobile and Ubiquitous Computing

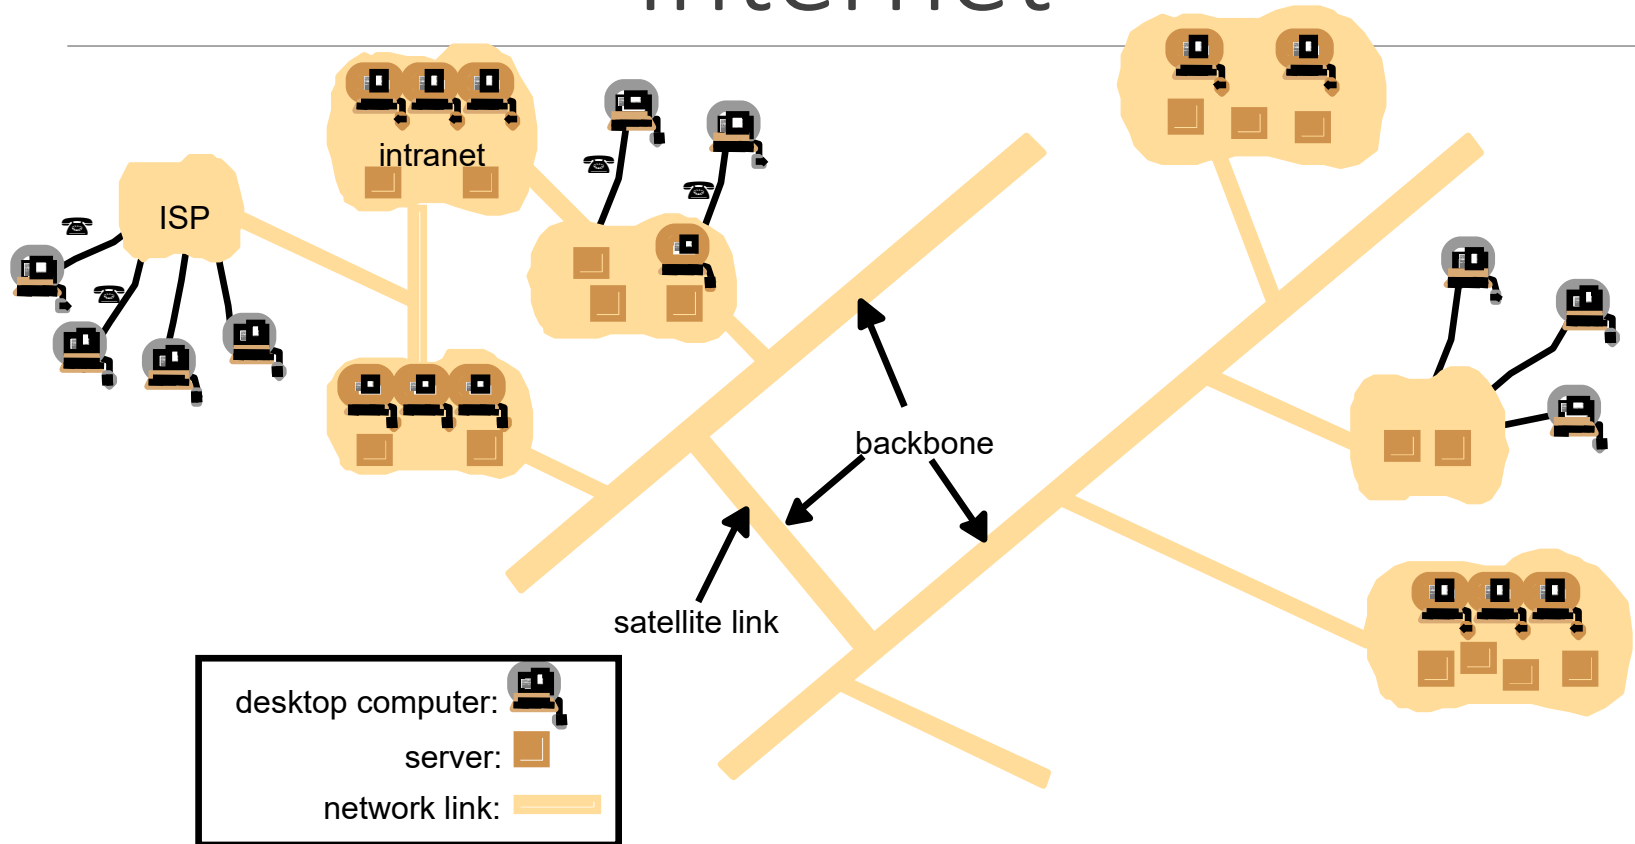
Database Management System



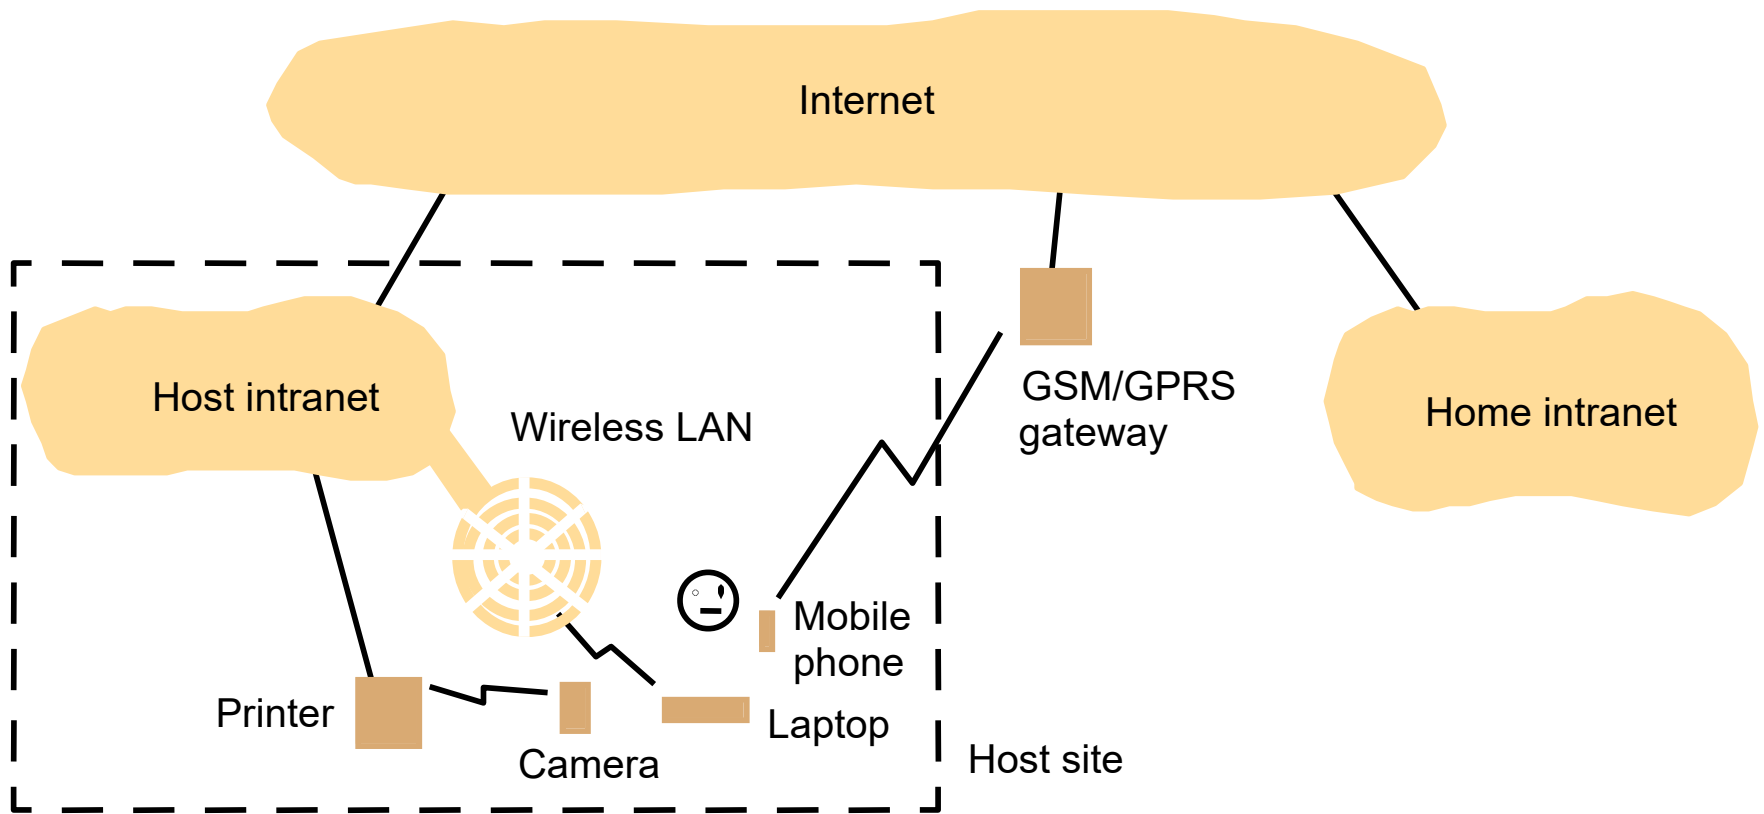
Automatic Teller Machine Network



Internet



Mobile and Ubiquitous Computing



Why Distributed Computing?

- Inherent distribution
 - ❖ connects customers, suppliers, and companies at different sites.
- Speedup - improved performance
- Fault tolerance – handling errors
- Resource Sharing
 - ❖ Exploitation of special hardware
- Scalability
- Flexibility

Why are Distributed Systems Hard?

- Scale
 - ❖ numeric, geographic, administrative
- Loss of control over parts of the system
- Unreliability of message passing
 - ❖ unreliable communication, insecure communication, costly communication
- Failure
 - ❖ Parts of the system are down or inaccessible
 - ❖ Independent failure is desirable

Flynn's taxonomy for computer's HW

Computer HW

```
graph TD; A[Computer HW] --- B[SISC  
Single Instruction Single Data  
(Ordinary PC)]; A --- C[SIMD  
Single Instruction Multiple Data  
(Array processor)]; A --- D[MISD  
Multiple Instruction Single Data  
(Pipelined processors)]; A --- E[MIMD  
Multiple Instruction Multiple Data  
(Multiprocessors or multicomputers)];
```

SISC

Single Instruction Single Data
(Ordinary PC)

SIMD

Single Instruction Multiple Data
(Array processor)

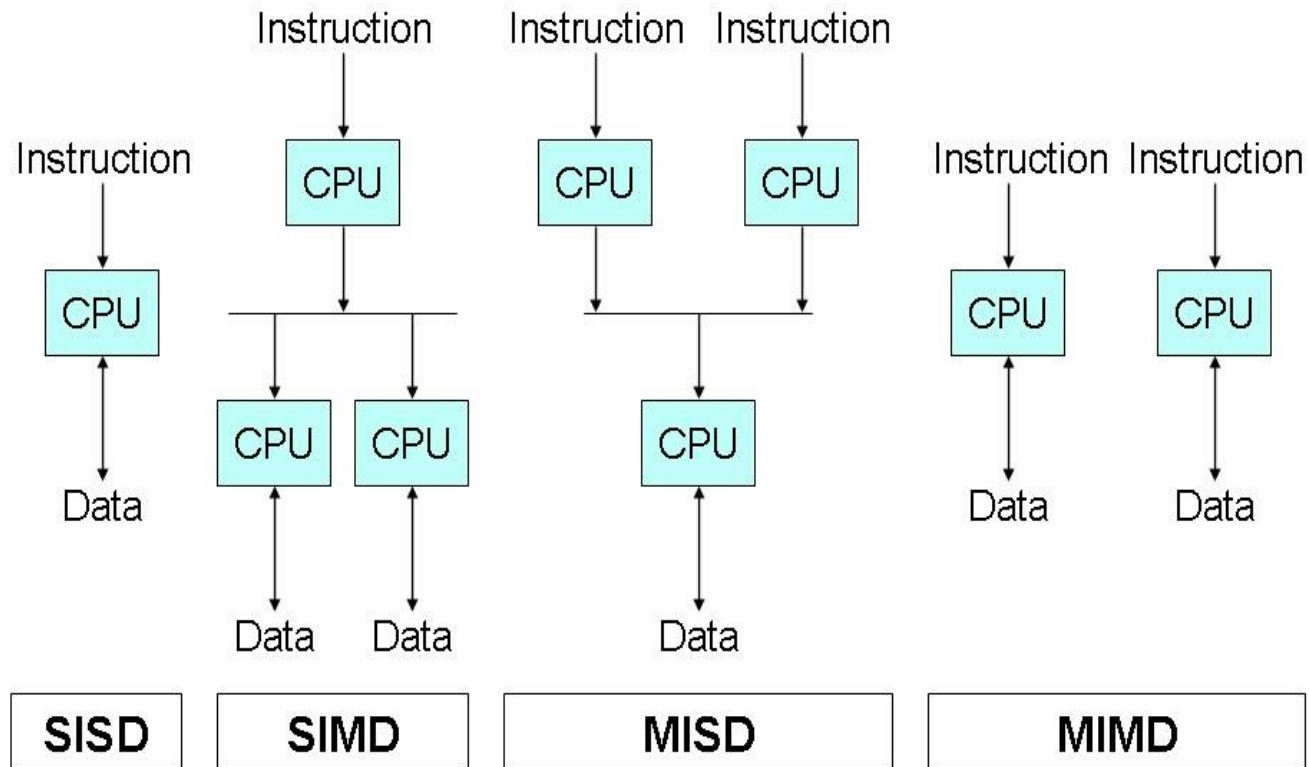
MISD

Multiple Instruction Single Data
(Pipelined processors)

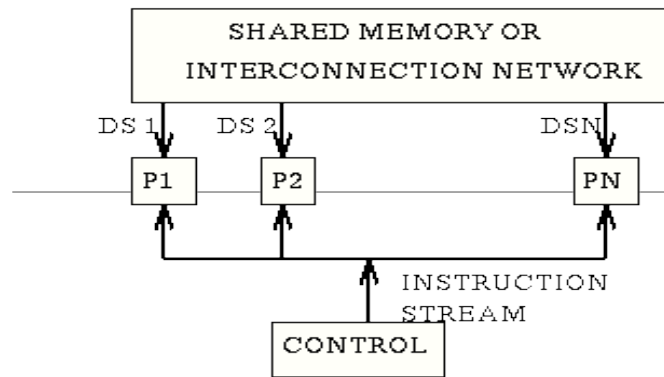
MIMD

Multiple Instruction Multiple Data
(Multiprocessors or multicomputers)

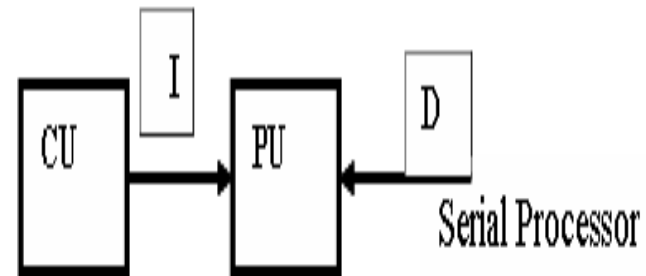
Flynn's Taxonomy



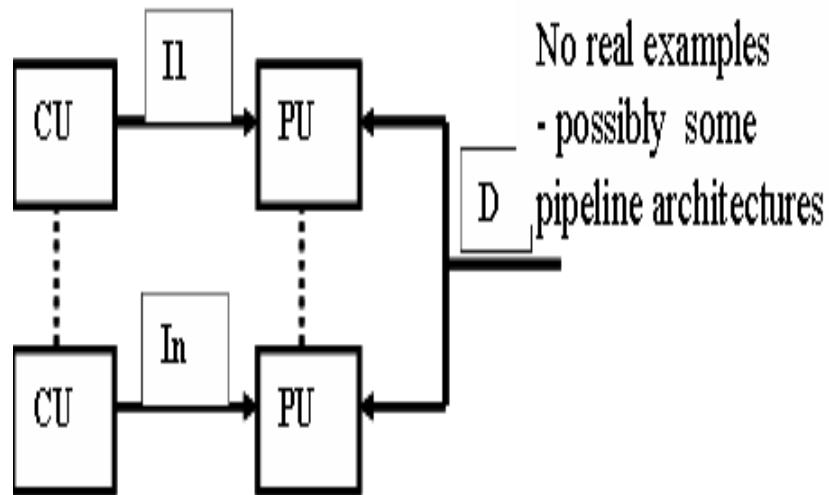
SIMD



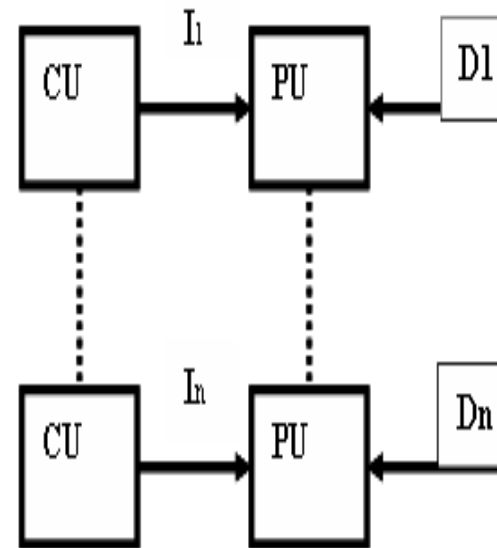
P = PROCESSOR
DS = DATA STREAM



SISD



MISD



MIMD

Multiprocessor
and
Multicomputer

(SISD) single instruction stream and a single data stream

- It has only one CPU and also called von Neumann computers, that starts from personal computers (PC) to large mainframes.
- A sequential machine process one instruction at a time on single data in single processing element.
- Communication are made between memory and CPU using Bus.
- Speed limited by CPU speed and bus speed.



(SISD) single instruction stream and a single data stream

- To compute the sum of N numbers a_1, a_2, \dots, a_N the processor needs to gain access to memory N consecutive times (to receive one number).
- $N-1$ additions are executed in sequence. Therefore the computation takes $O(N)$ operations. i.e. algorithms for SISD computers do not contain any parallelism, there is only one processor
- To speed up, add concurrency by:
 - 1- Concurrent execution of several different user's programs (multiprogramming). More than program share the CPU: one program uses CPU for some time and when it releases the CPU another program can use it (for example, round robin time scheduling technique).

- 2- Execute I/O operations simultaneously with the execution of user's program (do you know what is this called???).
- 3- Using parallel functional units inside the processing unit to be multifunction PU. Each unit will be assigned a specific function and has its own data stream so as more than one unit can function at the same time.
- 4- pipelining where the instruction or operation is broken into its element parts and each part is assigned to a processing element. Thus, phases of the instruction are now pipelined including fetching, decoding, operand fetch, ALU execution.

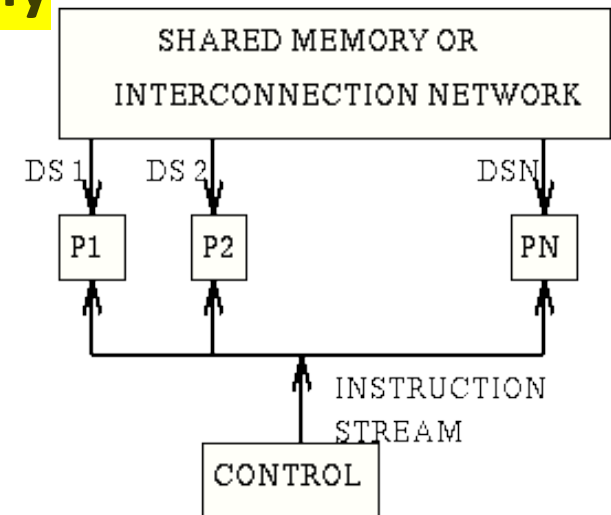
To summarize:

- *Sequential programming* refers to a set of ordered instructions executed one at a time on one CPU.
- *Concurrent programming* handles several operations at one time and doesn't require hardware support (using either one or multiple cores).
- *Parallel programming* executes multiple operations at the same time on multiple CPUs. All parallel programs are concurrent, running simultaneously, but not all concurrency is parallel. The reason is that parallelism is achievable only on multi-core devices.
- *Multitasking* concurrently performs multiple threads from different processes. Multitasking doesn't necessarily mean parallel execution, which is achieved only when using multiple CPUs.
- *Multithreading* extends the idea of multitasking; it's a form of concurrency that uses multiple, independent threads of execution from the same process. Each thread can run concurrently or in parallel, depending on the hardware support.

SIMD

Single Instruction Multiple Data (Array and vector processor)

- All N identical processors operate under the control of a single instruction stream issued by a central control unit.
- There are N data streams, one per processor so different data can be used in each processor.
- The processors operate **synchronously** and a **global clock** is used to ensure **lockstep** operation. i.e. at each step (global clock tick) all processors execute the same instruction, each on a different data stream.



P = PROCESSOR
DS = DATA STREAM

SIMD

Single Instruction Multiple Data (Array and vector processor)

- It is best used in applications that contains operation on matrices or vectors.
- Each processor will execute the same operation but on different data set which is usually a part of the matrix.
- The total result will be stored in the shared memory.
- The processing elements (PEs) are slaves under control of a master CU.
- The CU has many functions like: decoding the instruction, generating micro-operations needed for execution for the PEs, generating and broadcast common memory addresses and receiving the data from PEs and producing the final

results

- Array processors such as the ICL DAP (Distributed Array Processor) and vector computers such as the CRAY 1 & 2 and CYBER 205 fit into the SIMD category.
- For Example, Adding two matrices $A + B = C$.
Say we have two matrices A and B of order 2 and we have 4 processors.
- $A_{11} + B_{11} = C_{11} \dots A_{12} + B_{12} = C_{12}$
- $A_{21} + B_{21} = C_{21} \dots A_{22} + B_{22} = C_{22}$
- The same instruction is issued to all 4 processors (add the two numbers) and all processors execute the instructions simultaneously.
It takes one step as opposed to four steps on a sequential machine.
- An instruction could be a simple one (eg adding two numbers) or a complex one (eg merging two lists of numbers). Similarly the datum may be simple (one number) or complex (several numbers).

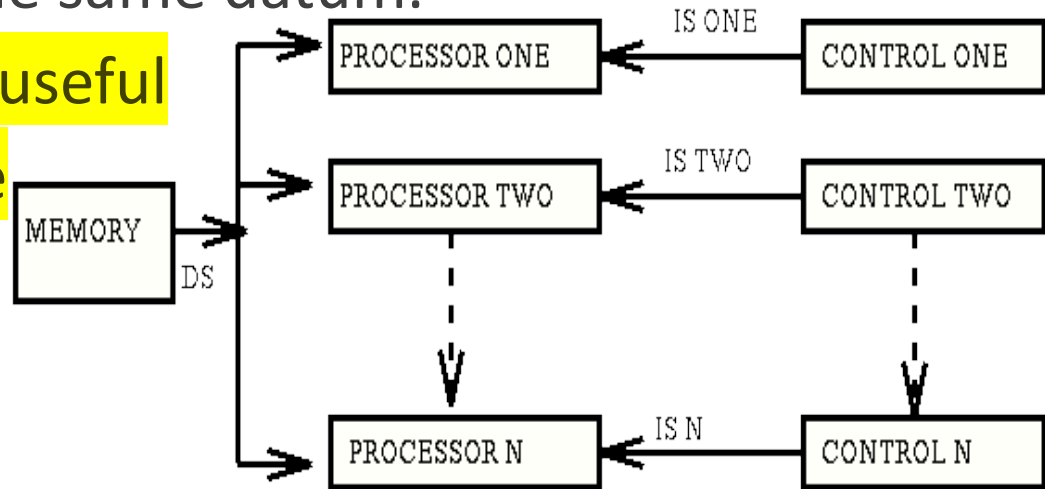
- Sometimes it may be necessary to have only a subset of the processors execute an instruction i.e. only some data needs to be operated on for that instruction.
-
- This information can be encoded in the instruction itself indicating whether
 - the processor is **active** (execute the instruction)
 - the processor is **inactive** (wait for the next instruction)
 - In most problems to be solved on SIMD (and MIMD) computers it is useful for the processors to be able to communicate with each other to exchange data or results. This can be done in two ways: by using a **shared memory and shared variables** or **some form of interconnection network** and message passing (distributed memory).

SDMI

Single Data Multiple Instruction (Pipelined processors)

- N processors, each with its own control unit, share a common memory.
- There are N streams of instructions (algorithms / programs) and **one** stream of data. Parallelism is achieved by letting the processors do different things at the same time on the same datum.

➤ MISD machines are useful in computations where the same input is to be subjected to several different operations.



IS = INSTRUCTION STREAM

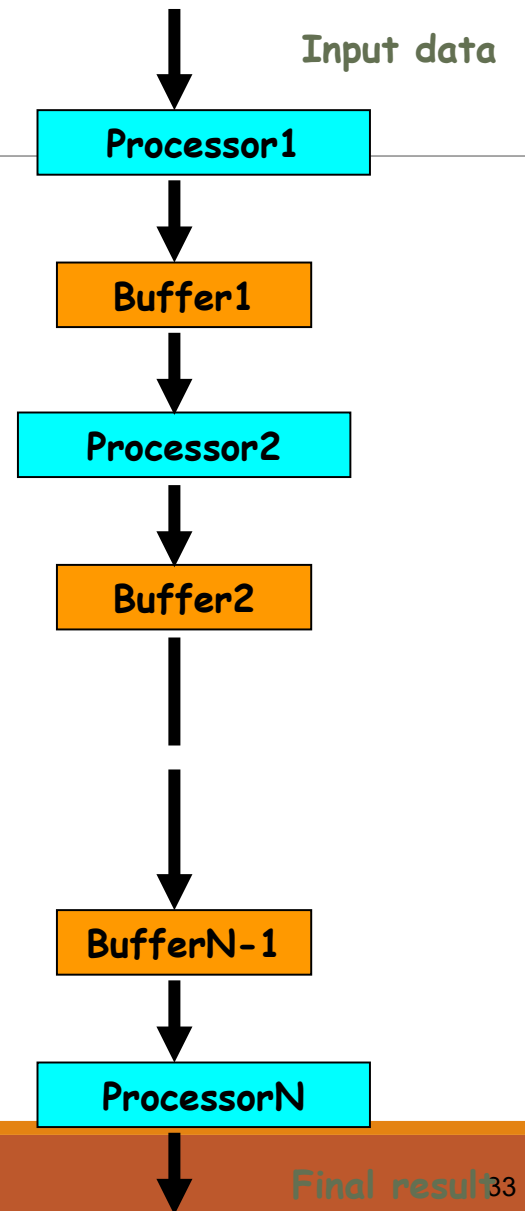
DS = DATA STREAM

MISD

Multiple Instruction Single Data (Pipelined processors)

- For example, **Checking whether a number Z is prime.**
A simple solution is to try all possible divisions of Z . Assume the number of processors, N , is given by $N = Z-2$. All processors take Z as input and tries to divide it by its associated divisor. So it is possible in **one step** to check if Z is prime. More realistically if $N < Z-2$ then a subset of divisors would be assigned to each processor.
- For most applications MISD is no commercial machines exist with this design.
- However, a pipelined processors is another example of MISD.

- Its basic idea is to break a complex time consuming function into a series of simpler sequential operations.
- The pipelined processors is then a machine that overlaps computations by subdividing and interleaving the operation of these sub-computations each on individual processor as happens in an assembly line fashion.
- The output of one stage is the input of the next.
- Buffers are used to hold output data temporally until the next stage finishes its data execution.
- Draw time overlap for different data streams

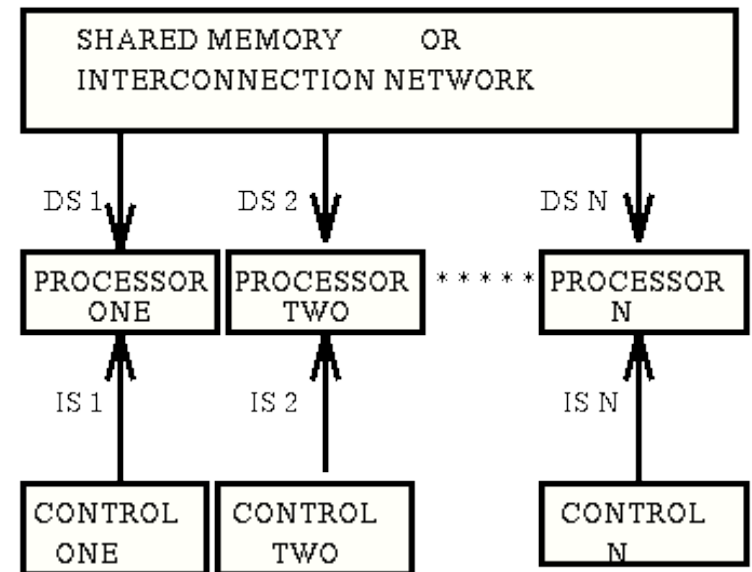


- Speedup = $NM/N+M-1$
- Speedup of a pipeline is defined to be: the number of times faster a pipeline component operates than a nonpipelined component that performs the same operation.
- N = no. of data streams
- M = Number of stages where each stage will require 1 time unit
- NM = total time taken to complete the N operations in nonpipelined component.
- For perfect usage of pipeline idea, time should be the same for all stages.
- If one stage takes less time, the remaining time will be wasted ??? How???

MIMD

Multiple Instruction Multiple Data (Multiprocessors or multicomputers)

- This is the most general and most powerful of our classification. We have N processors, N streams of instructions and N streams of data.
- Each processor is fully programmable and can execute its own program.
- It can operate under the control of an instruction stream issued by its own control unit (i.e. each processor is capable of executing its own program on a different data.
- This means that the processors operate **asynchronously** (typically) i.e. can be doing different things on different data at the same time.



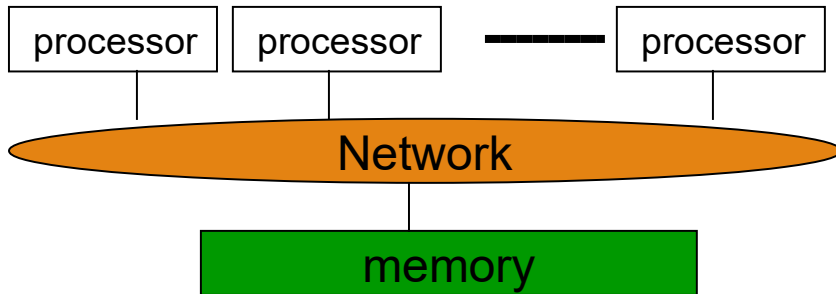
MIMD

Multiple Instruction Multiple Data (Multiprocessors or multicomputers)

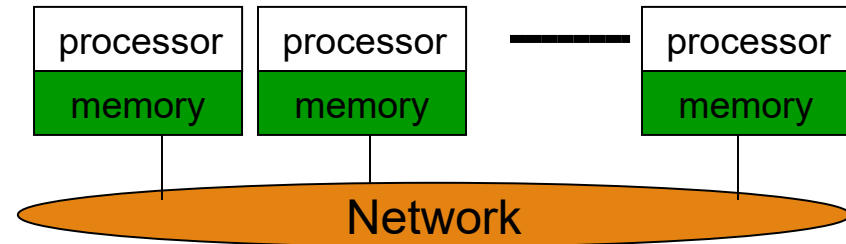
- As with SIMD computers, communication of data or results between processors can be via a shared memory or interconnection network.
- MIMD computers with shared memory are known as **multiprocessors** or **tightly coupled machines**. Examples are ENCORE, MULTIMAX, SEQUENT & BALANCE.
- MIMD computers with an interconnection network are known as **multicomputers** or **loosely coupled machines**. Examples are INTEL iPSC, NCUBE/7 and transputer networks.
- **Note**
Multicomputers are sometimes referred to as **distributed systems** and sometimes not. Some text books in the literature claim that Distributed systems should work together to accomplish **one** job.

Distributed Hardware

Physically shared memory



Physically distributed memory

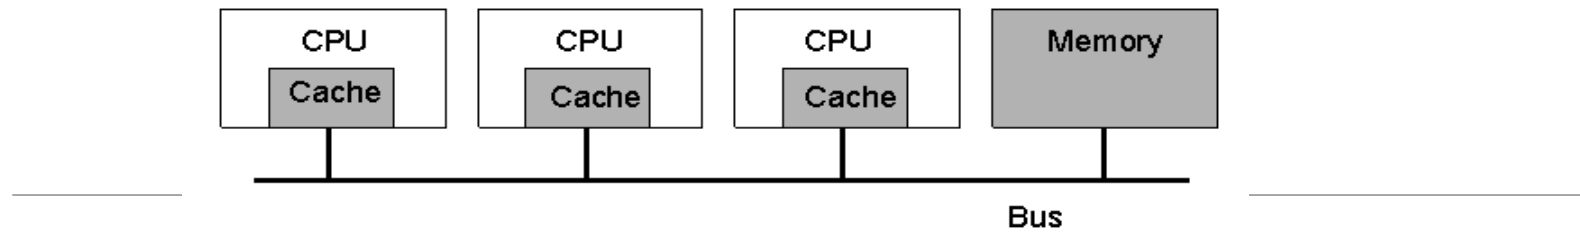


Physically shared/distributed memory and logically
shared/distributed memory

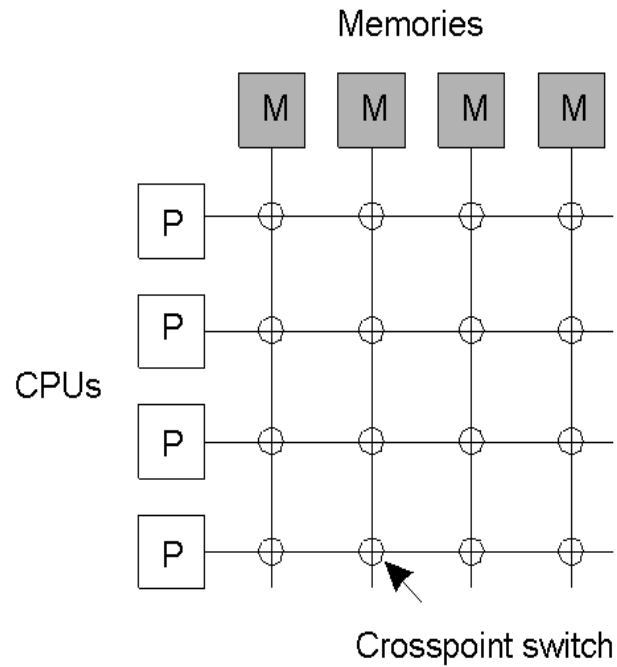
| | Logically shared | Logically distributed |
|------------------------|---------------------------|---------------------------|
| Physically shared | Common memory | Simulated message passing |
| Physically distributed | Distributed shared memory | message passing |

Tightly and loosely coupled systems

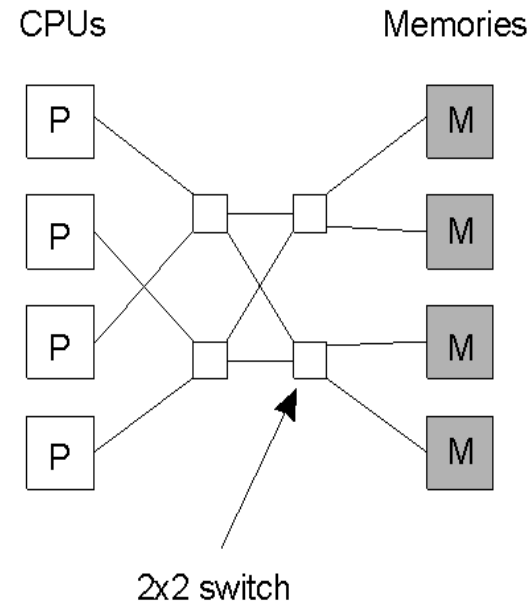
- Another dimension to our taxonomy is that in some systems the machines are tightly coupled and in others they are loosely coupled.
- In a tightly-coupled system, the delay experienced, when a message is sent from one computer to another is small, and the data rate is high that is, the number of bits per second that can be transferred is large.
- Typically, tightly coupled systems are referred to as multiprocessors when a number of processors each having its memory and are organized on a single chassis and have access to common memory through high speed bus under control of single OS. Examples of this are: the hypercube and the tree network of processors.
- There must be strong interaction at HW and OS levels, at the dataset level, the program execution level.



.A bus-based multiprocessor

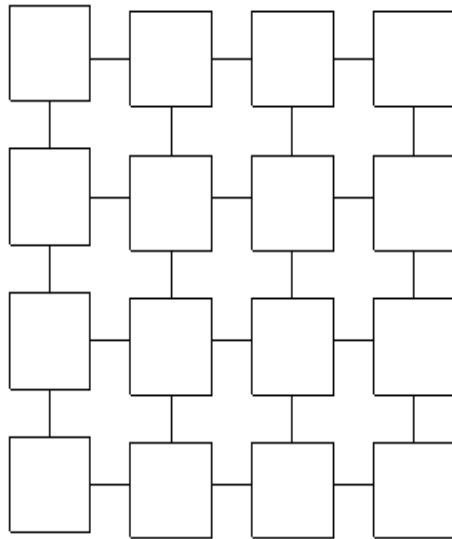


(a)

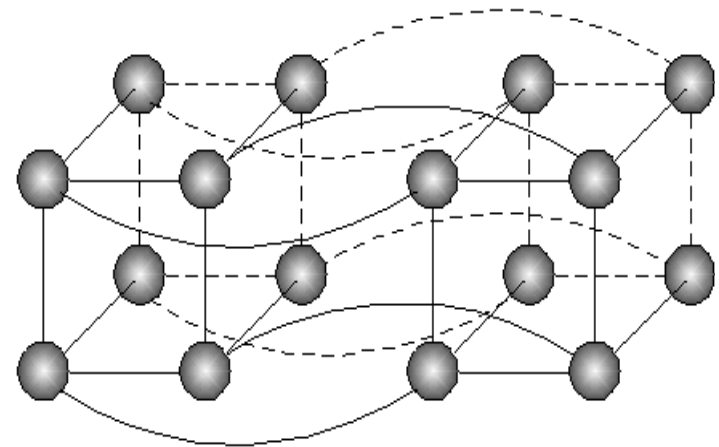


(b)

.A crossbar switch. (b) An omega switching network (a)



(a)



(b)

.Grid multiprocessor (b) Hypercube multiprocessor (a)

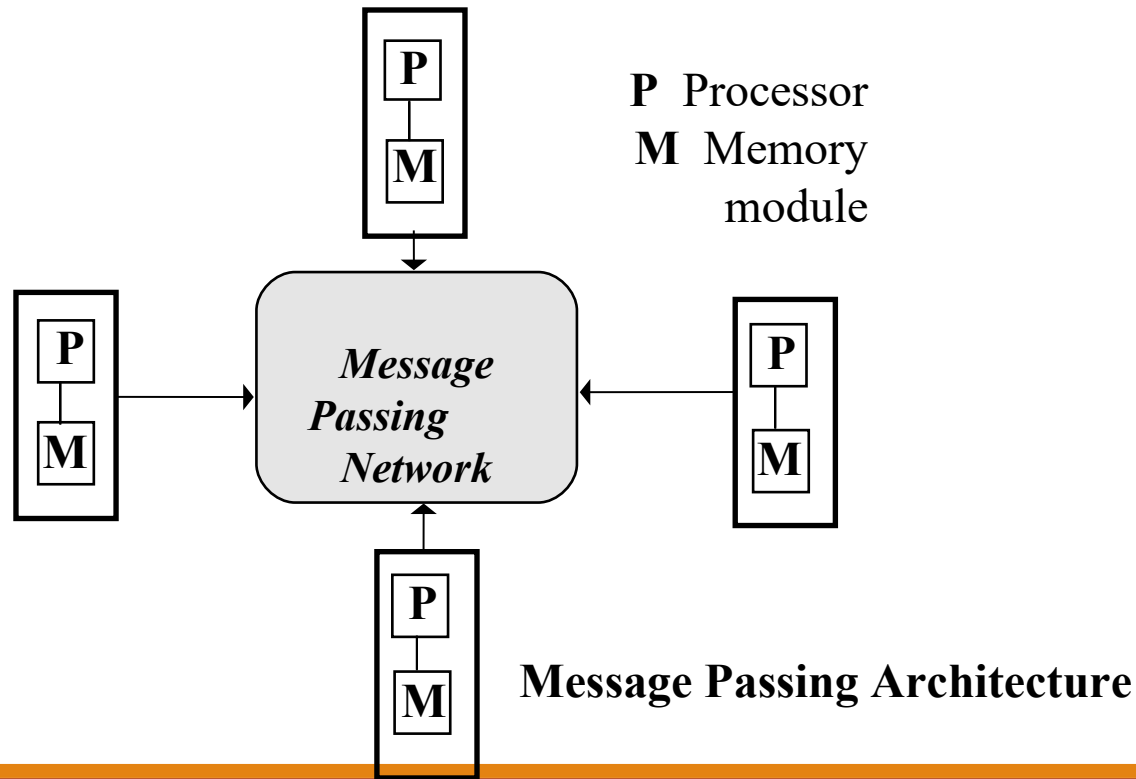
➤ Grid Microprocessor

- rectangular array of CPUs with connections only to neighbors
- messages from one corner to the opposite corner in an n by n grid require $(n-1)+(n-1)$ hops

➤ Hypercube Microprocessor

- an n -dimensional cube with CPUs at the vertices
- there are 2^n CPUs
- messages from one CPU to a CPU most distant from it require n hops ($= \log_2(\text{number of CPUs})$)

- In a loosely-coupled system, there are multiple computers (network) that communicate at the I/O level through intermachine message, thus the delay is large and the data rate is low.
- In order to complete a transfer of data between two processors, a cooperation is required between the two processors at the OS level.



- Tightly-coupled systems tend to be used more as parallel systems (working on a single problem) and loosely-coupled ones tend to be used as distributed systems (working on many unrelated problems), although this is not always true.
- It is possible to have a system that is physically loosely coupled but logically tightly coupled. This could be made using OS constructs.
- Grid and clusters are examples of loosely coupled systems that can work in an orderly fashion to accomplish one job.
- When working to solve a certain problem, these processors should be organized using single operating system that organizes the interaction between them, assigns tasks to each processor and to ensure the synchronization of the assigned tasks.

Table ■ Computer Generations in 50 Years

| Generation Period | Technology and Architecture | Software and Operating System | Representative Systems |
|--------------------------------|--|---|--|
| First 1946-1956 | Vacuum tubes and relay memory, single-bit CPU with accumulator-based instruction set | Machine/assembly languages, programs without subroutines | ENIAC, IBM 701, Princeton IAS |
| Second 1956-1967 | Discrete transistors, core memory, floating-point accelerator, I/O channels | Algol and Fortran with compilers, batch processing OS | IBM 7030, CDC 1604, Univac LARC |
| Third 1967-1978 | Integrated circuits, pipelined CPU, microprogrammed control unit | C language, multiprogramming, timesharing OS | PDP-11 IBM 360/370, CDC 6600 |
| Fourth 1978-1989 | VLSI microprocessors, solid-state memory, multiprocessors, vector supercomputers | Symmetric multiprocessing, parallelizing compilers, message-passing libraries | IBM PC, VAX 9000, Cray X/MP |
| Fifth 1990- present | ULSI circuits, scalable parallel computers, workstation clusters, Intranet, Internet | Java, microkernels, Multithreading, distributed OS, World-Wide Web | IBM SP2, SGI Origin 2000, Digital TruCluster |

In hardware technology, the first generation used vacuum tubes and relay memory. The second generation was marked by the use of discrete transistors and core memory.

Similar terms for distributed systems

- Parallel systems
- Concurrent systems
- Networked systems
- Decentralized systems

Parallel computing

- *A type of computation in which many calculations or the execution of processes are carried out simultaneously.*
- Large problems can often be divided into smaller ones, which can then be solved at the same time.
- There are several different forms of parallel computing: *bit-level, instruction-level, data, and task parallelism.*
- In parallel computing, a computational task is typically broken down in several, often many, very similar subtasks that can be processed independently and whose results are combined afterwards, upon completion.

Concurrent Systems

- In computer science, **concurrency** refers to the ability of different parts or units of a program, algorithm, or problem to be executed out-of-order or in partial order, without affecting the final outcome.
- This allows for parallel execution of the concurrent units, which can significantly improve overall speed of the execution in multi-processor and multi-core systems.
- In more technical terms, concurrency refers to the decomposability property of a program, algorithm, or problem into order-independent or partially-ordered components or units.

Concurrent Systems

- Concurrent systems such as **Operating systems and Database management systems** are generally designed to operate indefinitely, including automatic recovery from failure, and not terminate unexpectedly.
- Because **they use shared resources**, concurrent systems in general require the inclusion of some kind of **arbiter** somewhere in their implementation.

Parallel computing

- Parallelism has been employed for many years, mainly in high-performance computing, but interest in it has grown lately due to the physical constraints preventing frequency scaling.
- As power consumption (and consequently heat generation) by computers has become a concern in recent years, parallel computing has become the dominant paradigm in computer architecture, mainly in the form of *multi-core processors*.
- Parallel computing is closely related to concurrent computing—they are frequently used together, and often conflated.

Parallel and Concurrent

- Though the two are distinct:
- it is possible to have parallelism without concurrency (such as bit-level parallelism),
- and concurrency without parallelism (such as multitasking by time-sharing on a single-core CPU).
- In contrast to parallel computing, in concurrent computing, the various processes often do not address related tasks; when they do, as is typical in distributed computing, the separate tasks may have a varied nature and often require some inter-process communication during execution.

Parallel computers can be roughly classified

- According to the level at which the hardware supports parallelism, with multi-core and multi-processor computers having multiple processing elements within a single machine,
- While clusters, Cloud, and grids use multiple computers to work on the same task.
- Specialized parallel computer architectures are sometimes used alongside traditional processors, for accelerating specific tasks.
- In some cases parallelism is transparent to the programmer, such as in bit-level or instruction-level parallelism, but explicitly *parallel algorithms*, particularly those that use concurrency, are more difficult to write than sequential ones, because concurrency introduces several new classes of potential software bugs.

CENTRALIZED SYSTEMS

- Centralized systems are systems that use **client/server** architecture where one or more client nodes are directly connected to a central server.
- This is the most commonly used type of system in many organizations where a client sends a request to a company server and receives the response.
- **Scaling –**
Only vertical scaling on the central server is possible. Horizontal scaling will contradict the single central unit characteristic of this system of a single central entity.

Components of Centralized System

- Node (Computer, Mobile, etc.).
 - Server.
 - Communication link (Cables, Wi-Fi, etc.).
-

Limitations of Centralized System –

- Can't scale up vertically after a certain limit – After a limit, even if you increase the hardware and software capabilities of the server node, the performance will not increase appreciably leading to a cost/benefit ratio < 1 .
- Bottlenecks can appear when the traffic spikes – as the server can only have a finite number of open ports to which can listen to connections from client nodes. So, when high traffic occurs like a shopping sale, the server can essentially suffer a Denial-of-Service attack or Distributed Denial-of-Service attack.

Characteristics of Centralized System

- **Presence of a global clock:** As the entire system consists of a central node(a server/ a master) and many client nodes(a computer/ a slave), all client nodes sync up with the global clock(the clock of the central node).
- **One single central unit:** One single central unit which serves/coordinates all the other nodes in the system.
- **Dependent failure of components:** Central node failure causes the entire system to fail. This makes sense because when the server is down, no other entity is there to send/receive responses/requests.

Use Cases

- Centralized databases – all the data in one server for use.
- Single-player games like Need For Speed, GTA Vice City – an entire game in one system (commonly, a Personal Computer)
- Application development by deploying test servers leading to easy debugging, easy deployment, easy simulation
- Personal Computers

Decentralized systems

- In decentralized systems, every node makes its own decision.
- The final behavior of the system is the aggregate of the decisions of the individual nodes.
- Note that there is no single entity that receives and responds to the request.

Characteristics of Decentralized System

- **Lack of a global clock:** Every node is independent of each other and hence, has different clocks that they run and follow.
- **Multiple central units (Computers/Nodes/Servers):** More than one central unit which can listen for connections from other nodes
- **Dependent failure of components:** one central node failure causes a part of the system to fail; not the whole system

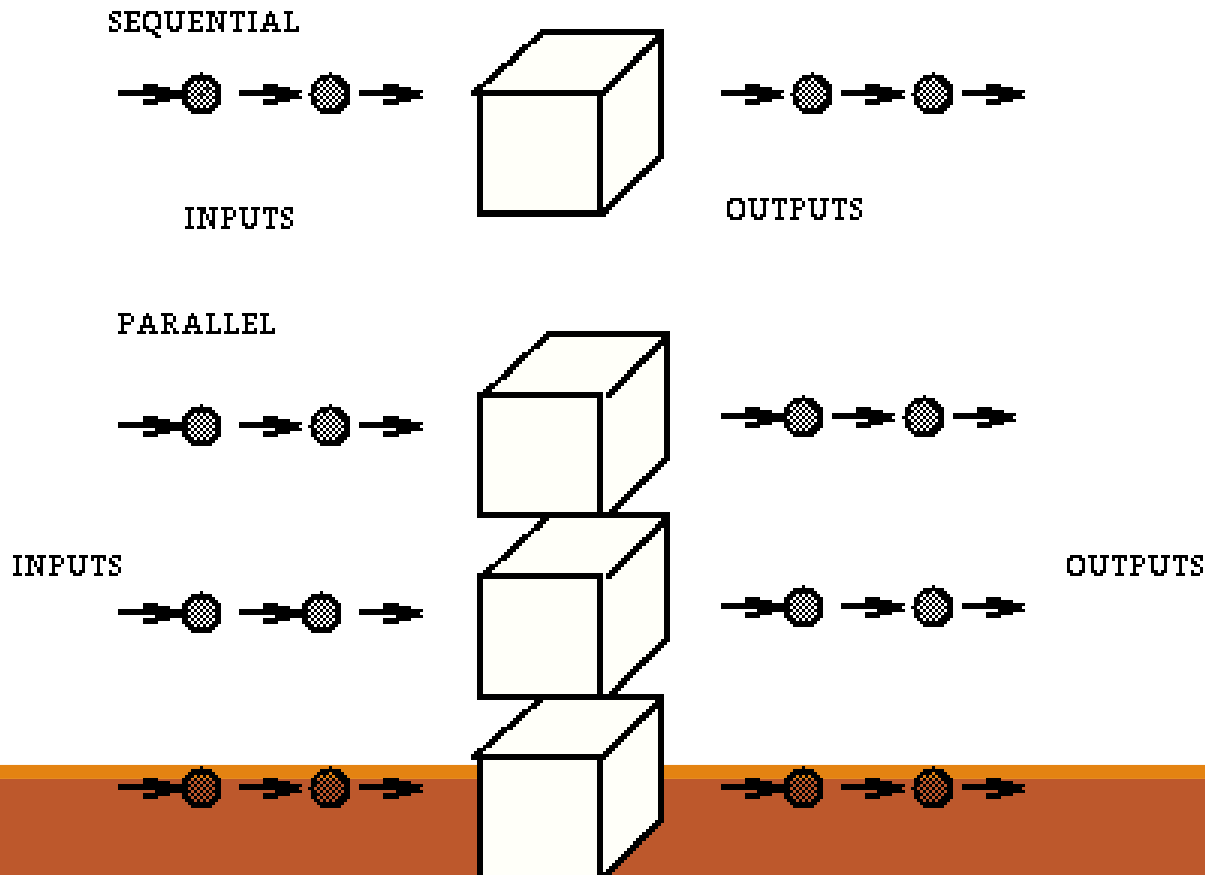
Characteristics of Decentralized System

- **Lack of a global clock:** Every node is independent of each other and hence, has different clocks that they run and follow.
- **Multiple central units (Computers/Nodes/Servers):** More than one central unit which can listen for connections from other nodes
- **Dependent failure of components:** one central node failure causes a part of the system to fail; not the whole system

The Need For High Performance Computers

- Many of today's applications such as weather prediction, aerodynamics and artificial intelligence are very computationally intensive and require vast amounts of processing power.
- To calculate a 24 hour weather forecast for the many countries requires about 10^{12} operations to be performed. This would take about 2.7 hours on a Cray-1 (capable of 10^8 operations per second).
- So to give accurate long range forecasts (e.g. a week) much more powerful computers are needed.
- One way of doing this is to use faster electronic components. The limiting factor is however the speed of light.
- The speed of light is $3 * 10^8$ m/s.
- Considering two electronic devices (each capable of performing 10^{12} operations per second) 0.5mm apart.
- It takes longer for a signal to travel between them than it takes for either of them to process it (10^{-12} seconds).
- So producing faster components is ultimately of no good.

- So it appears that the only way forward is to use PARALLELISM.
- The idea here is that if several operations can be performed simultaneously then the total computation time is reduced.
- The distributed system version has the **potential** of being 3 times as fast as the sequential machine.



Defining Speedup and Efficiency

- A **Parallel Algorithm** is an algorithm for the execution of a program which involves the running of two or more processes on two or more processors simultaneously.
- Two important measures of the quality of parallel algorithms are **speedup** and **efficiency**.
- If T_s is the time taken to run the fastest serial algorithm on one processor and if T_p is the time taken by a parallel algorithm on N processors then

- **Speedup** = $S_N = T_s / T_p$ $\frac{T_{\text{Sequential}}}{T_{\text{Parallel}}} = S_N$

and the efficiency of the parallel algorithm is given by

- **Efficiency** = $E_N = S_N / N$ $\frac{T_s \times N}{T_p}$

Example on Speedup and Efficiency

- If the best known serial algorithm takes 8 seconds i.e. $T_s = 8$, while a parallel algorithm takes 2 seconds using 5 processors, then

$$SN = T_s / T_p = 8 / 2 = \mathbf{4} \text{ and}$$

$$EN = SN / N = 4 / 5 = \mathbf{0.8 = 80\%}$$

- i.e. the parallel algorithm exhibits a speedup of 4 with 5 processors giving an 80% efficiency.

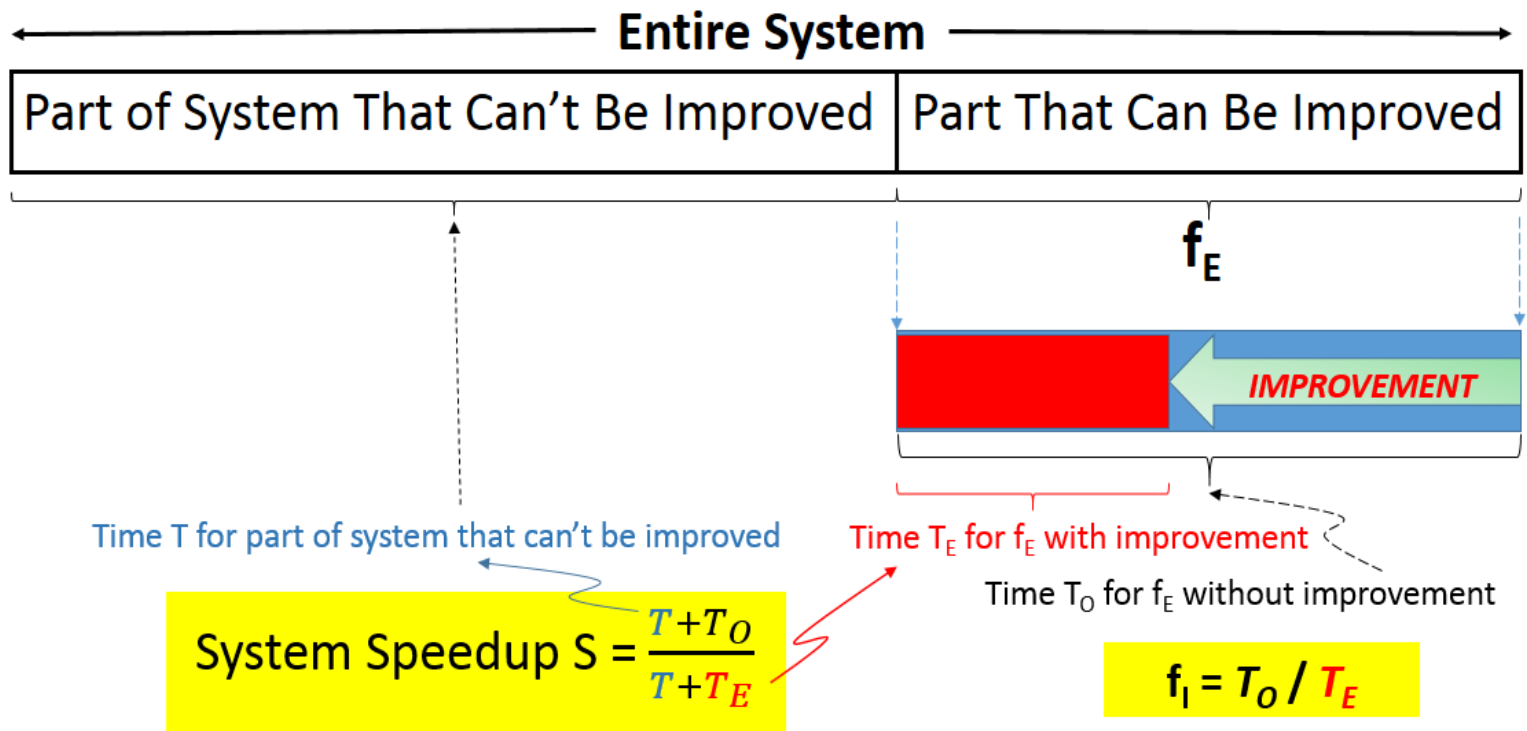
- A slightly different definition of speedup also exists.
- The time taken by the **parallel** algorithm on one processor divided by the time taken by the parallel algorithm on N processors.
- However, this is misleading since many parallel algorithms contain extra operations to accommodate the parallelism (e.g the communication) so the result is T_s is increased thus exaggerating the speedup.
- Which ever definition is used the ideal is to produce **linear speedup** i.e. produce a speedup of N using N processors and an efficiency of 1 (100%).
- However, in practice the speedup is reduced from its ideal value of N (the efficiency is bounded from above by 1).

Amdahl's law: Maximum speed-up when the algorithm is not 100% parallel

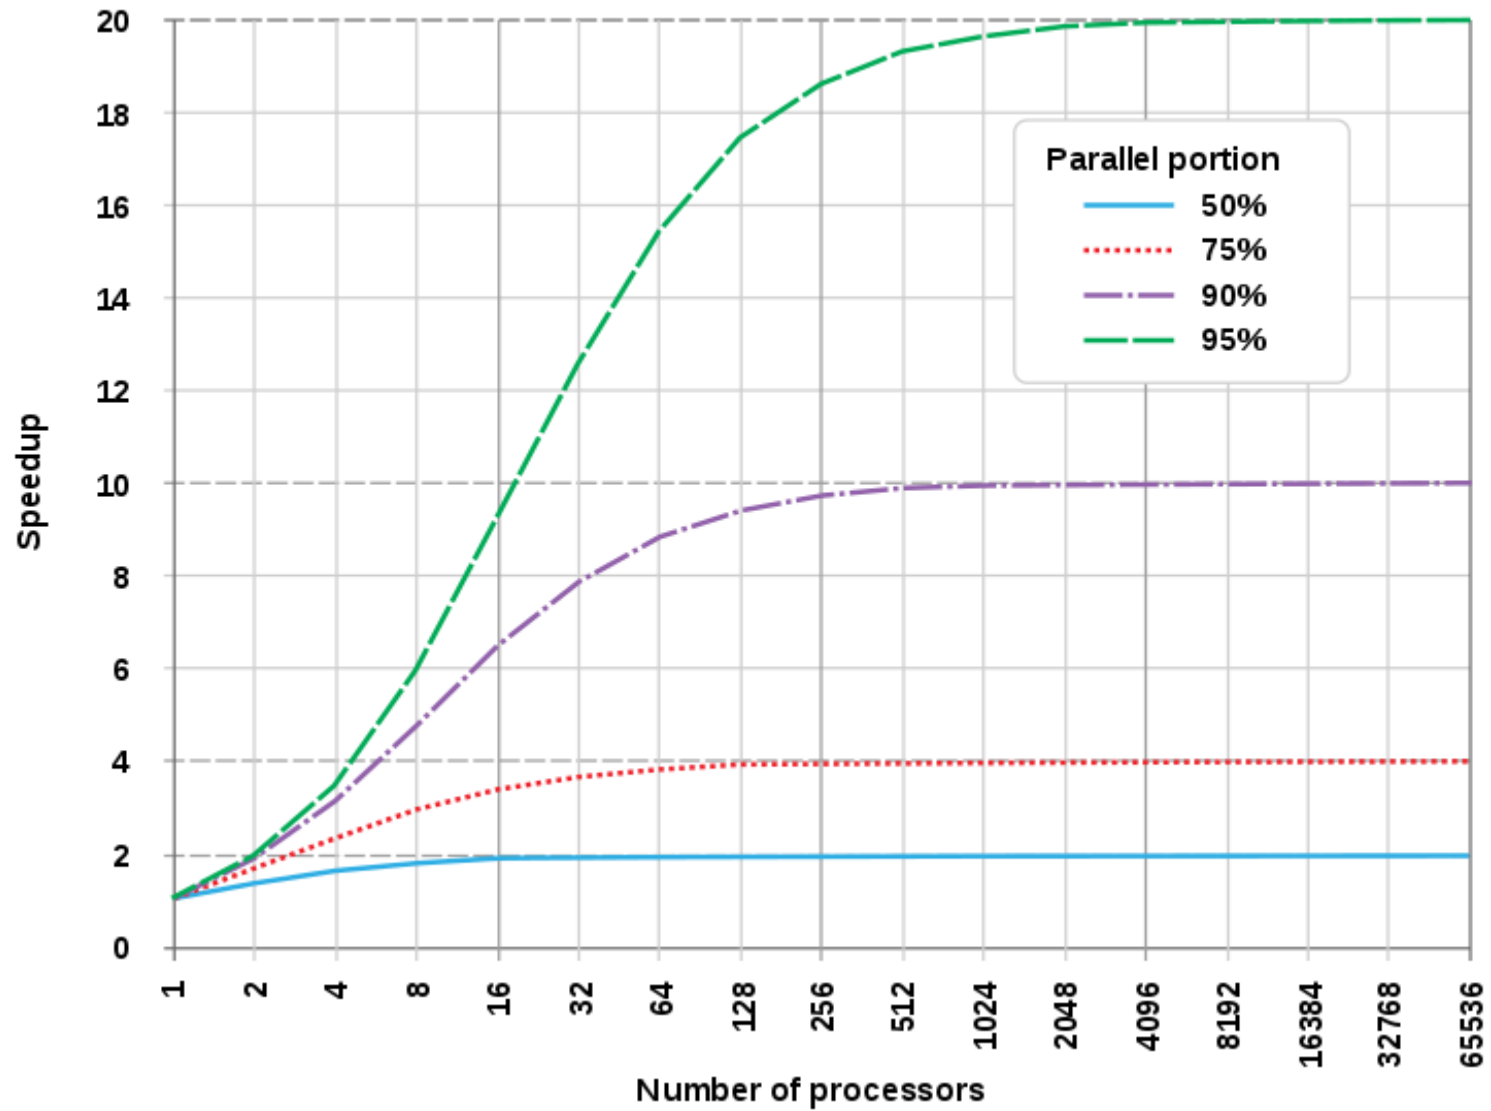
- Amdahl's law is a formula which gives the theoretical speedup in the execution at fixed **workload** that can be expected of a system whose resources (i.e. no of processors) are improved.
- If F is the fraction of a calculation that is sequential, and (1-F) is the fraction that can be parallelized,
- The maximum speed-up that can be achieved by using P processors is
Max speed-up = $1/(F+(1-F)/P)$
- Example:
- if 90% of a calculation can be parallelized (i.e. 10% is sequential) then the maximum speed-up which can be achieved on 5 processors is
Max speed-up = $1/(0.1+(1-0.1)/5)$ or roughly 3.6
- (i.e. the program can theoretically run 3.6 times faster on five processors than on one)

What Kinds of Problems Do We Solve With Amdahl's Law?

→ Performance Improvement Problems!!



Amdahl's Law



Amdahl's law

- In the previous example, If 90% of a calculation can be parallelized then the maximum speed-up on 10 processors is $1/(0.1+(1-0.1)/10)$ or 5.3 (i.e. investing twice as much hardware speeds the calculation up by about 50%).
- If 90% of a calculation can be parallelized then the maximum speed-up on 20 processors is $1/(0.1+(1-0.1)/20)$ or 6.9 (i.e. doubling the hardware again speeds up the calculation by only 30%).
- The point that Amdahl was trying to make was that using lots of parallel processors was not a viable way of achieving the sort of speed-ups that people were looking for.

Factors that limit speedup

1. Software Overhead

- Even with a completely equivalent algorithm, software overhead arises in the concurrent implementation. (e.g. there may be additional index calculations necessitated by the manner in which data are "split up" among processors.)
i.e. there is generally more lines of code to be executed in the parallel program than the sequential program.

2. Load Balancing

- Speedup is generally limited by the speed of the slowest node. So an important consideration is to ensure that each node performs the same amount of work. i.e. the system is load balanced.

3. Communication Overhead

- Assuming that communication and calculation cannot be overlapped, then any time spent communicating the data between processors directly degrades the speedup. (because the processors are not calculating).

Advantages of Distributed Systems over Centralized Systems

Performance:

- very often a collection of processors can provide higher performance (and better price/performance ratio) than a centralized computer. A distributed system may have more computing power than a centralized mainframe or supercomputer

Distribution:

- many applications involve, by their nature, spatially separated machines (banking, commercial, automotive system).

Reliability (fault tolerance):

- if some of the machines crash, the system can survive.

Incremental growth:

- as requirements on processing power grow, new machines can be added incrementally.

Sharing of data/resources:

- shared data is essential to many applications (banking, computer supported cooperative work, reservation systems); other resources can be also shared (e.g. expensive printers).

Communication:

- facilitates human-to-human communication.

Disadvantages of Distributed Systems

Difficulties of developing distributed software:

- How should operating systems, programming languages and applications look like?

Networking problems:

- several problems are created by the network infrastructure, which have to be dealt with: loss of messages, overloading, ...

Security problems:

- sharing generates the problem of data security.

Characteristics of Distributed System

- In summary we can conclude that the careful design and implementation of a distributed system must consider some or all of the following characteristics.
- Resource sharing
- Ability to use any hardware, software or data (called resources) anywhere in the system.
- Resource provided by a computer which is a member of a distributed system can be shared by clients and other members of the system via a network.
- “Resource Manager” is a software module (based on set of management policies) that provides interfaces which enables resource to be manipulated by clients.
- Provider and user interact with each other. (Client/server Model)

- Openness (extensible in various ways)
 - Openness is concerned with extensions and improvements of distributed systems.
-
- Detailed interfaces of components need to be published.
 - New components have to be integrated with existing components.
 - Differences in data representation of interface types on different processors (of different vendors) have to be resolved.
 - New resource sharing services can be incorporated without disruption or duplication of existing services.

➤ Transparency

- Hide all unnecessary details from users.
- Location transparency - clients do not need to know the location of the servers. Other types ????

➤ Shareability

- Allows the comprising systems to use each other's resources.

➤ Local Autonomy

- Each processor is able to manage its local resources.

➤ Improved reliability and availability

- Disruption would not stop the whole system from providing its services as resources Spread across multiple computers.

➤ Improved performance

- Load balancing, resource replication.
 - Combined processing power of multiple computers provides much more processing Power than a centralized system with multiple CPUs.
-
- Several factors are influencing the performance of a distributed system:
 - ❖ The performance of individual workstations.
 - ❖ The speed of the communication infrastructure.
 - ❖ Extent to which reliability (fault tolerance) is provided (replication and preservation of coherence imply large overheads).
 - ❖ Flexibility in workload allocation: for example, idle processors (workstations) could be allocated automatically to a user's task.

➤ Concurrency

- All concurrent access must be synchronized to avoid problems such as lost update, dirty read, incorrect summary & unrepeatable read.

➤ Scalability

- Flexible to grow in size, efficiently utilize the new s/w & h/w added.
- The system should remain efficient even with a significant increase in the number of users and resources connected:
 - cost of adding resources should be reasonable;
 - performance loss with increased number of users and resources should be controlled;

➤ Heterogeneity

- Not having the same things like; programming languages, OSs, HW platforms, network protocols.
- Distributed applications are typically heterogeneous:
 - different hardware: mainframes, workstations, PCs, servers, etc.;

- different software: UNIX, MSWindows, IBM OS/2, Real-time OSs, etc.;
 - unconventional devices: teller machines, telephone switches, robots, manufacturing systems, etc.;
 - diverse networks and protocols: Ethernet, FDDI, ATM, TCP/IP, Novell Netware, etc.
-
- The solution
 - *Middleware*, an additional software layer to mask heterogeneity
 - Fault tolerance
 - The system appropriately handles errors occurred.
 - The system continues to work even if parts of it fails.
 - Accomplished via HW redundancy, software recovery.
 - **Fault-tolerance** is a main issue related to reliability: the system has to detect faults and act in a reasonable way:

-
- ***mask the fault***: continue to work with possibly reduced performance but without loss of data/ information.
 - ***fail gracefully***: react to the fault in a predictable way and possibly stop functionality for a short period, but without loss of data/information.
 - Data on the system must not be lost, and copies stored redundantly on different servers must be kept **consistent**.
 - The more copies kept, the better the availability, but keeping consistency becomes more difficult.

Transparency

- How to achieve the single system image?
 - How to "fool" everyone into thinking that the collection
-

of machines is a "simple" computer?

- Access transparency
- local and remote resources are accessed using identical operations. ***Example: Navigation in the Web - SQL Queries***
- Location transparency
- - users cannot tell where hardware and software resources (CPUs, files, data bases) are located; the name of the resource shouldn't encode the location of the resource. ***Example: Pages in the Web - Tables in distributed databases***

- **Migration (mobility) transparency**

- resources should be free to move from one location to another without having their names changed.
 - ***Example: Web Pages***
-

- **Replication transparency**

- the system is free to make additional copies of files and other resources (for purpose of performance and/or reliability), without the users noticing.
- **Example: several copies of a file; at a certain request that copy is accessed which is the closest to the client.**

- **Concurrency transparency**

- - the users will not notice the existence of other users in the system (even if they access the same resources).
Example: Automatic teller machine network - Database management system

- **Failure transparency**
 - Enables the concealment of faults
 - applications should be able to complete their task despite failures occurring in certain components of the system. ***Example: Database Management System***
- **Performance transparency**
 - load variation should not lead to performance degradation. Allows the system to be reconfigured to
 - improve performance as loads vary.
 - This could be achieved by automatic reconfiguration as response to changes of the load; it is difficult to achieve
- **Scaling Transparency**
 - Allows the system and applications to expand in scale without change to the system structure or the application algorithms. ***Example: World-Wide-Web - Distributed Database***